



CATX

A JAX implementation of the *"Efficient Contextual Bandits with Continuous Actions"* paper

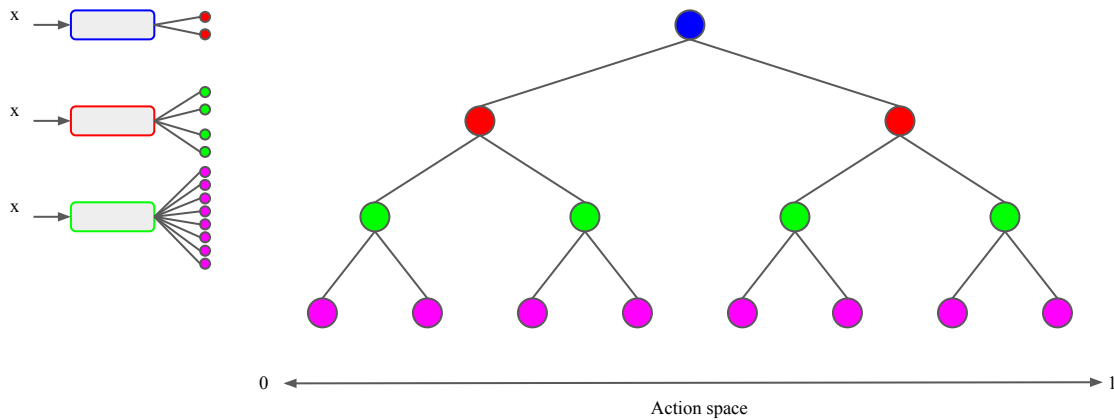
Tree

Tree

This example uses a tree of depth 3

At each depth there is neural network (depth 0: blue, depth 1: red, and depth 2: green)

Each neural network output layer dimension is $2^{(\text{depth}+1)}$



```
class Tree(hk.Module):
    def __init__(
        self,
        network_builder: NetworkBuilder,
        tree_params: TreeParameters,
        name: Optional[str] = None,
    ):

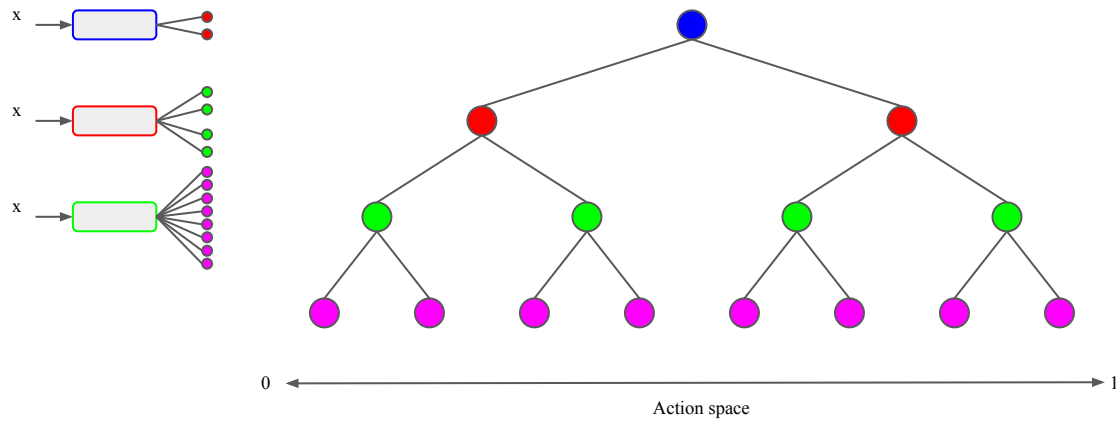
```

Tree parameters

Tree parameters


action space

```
@dataclass  
class TreeParameters:
```




Tree parameters

action spaces: each discretized action centroid covers $2 \times \text{bandwidth}$ of the action space




Tree parameters

action spaces: each discretized action centroid covers $2 \times \text{bandwidth}$ of the action space




Tree parameters

action spaces: each discretized action centroid covers $2 \times \text{bandwidth}$ of the action space



Tree parameters

action spaces: each discretized action centroid covers $2 \times \text{bandwidth}$ of the action space




Action query

```
def sample(  
    self, obs: Observations, epsilon: float  
) -> Tuple[Actions, Probabilities]:
```

Action query


forward pass of the tree by following the max of the logits



```
def sample(  
    self, obs: Observations, epsilon: float  
) -> Tuple[Actions, Probabilities]:
```

Action query

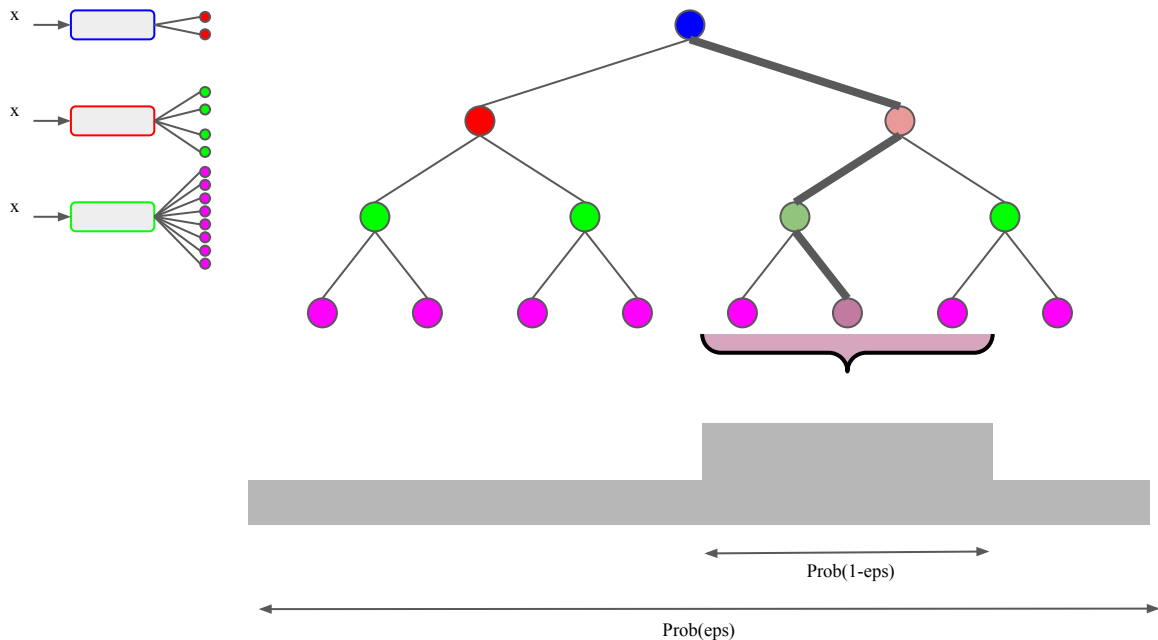
sample an action with eps-greedy



```
def sample(  
    self, obs: Observations, epsilon: float  
) -> Tuple[Actions, Probabilities]:
```

Action query

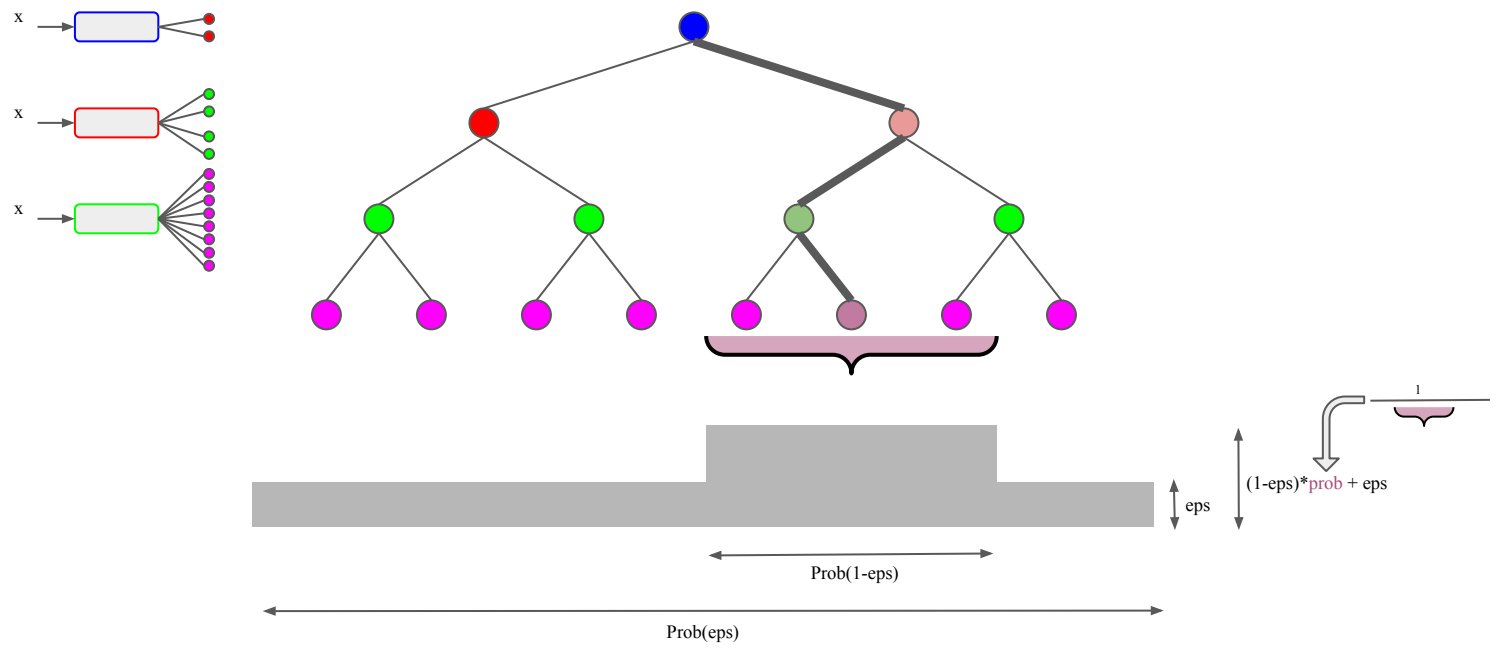
sample an action with eps-greedy



```
def sample(  
    self, obs: Observations, epsilon: float  
) -> Tuple[Actions, Probabilities]:
```

Action query

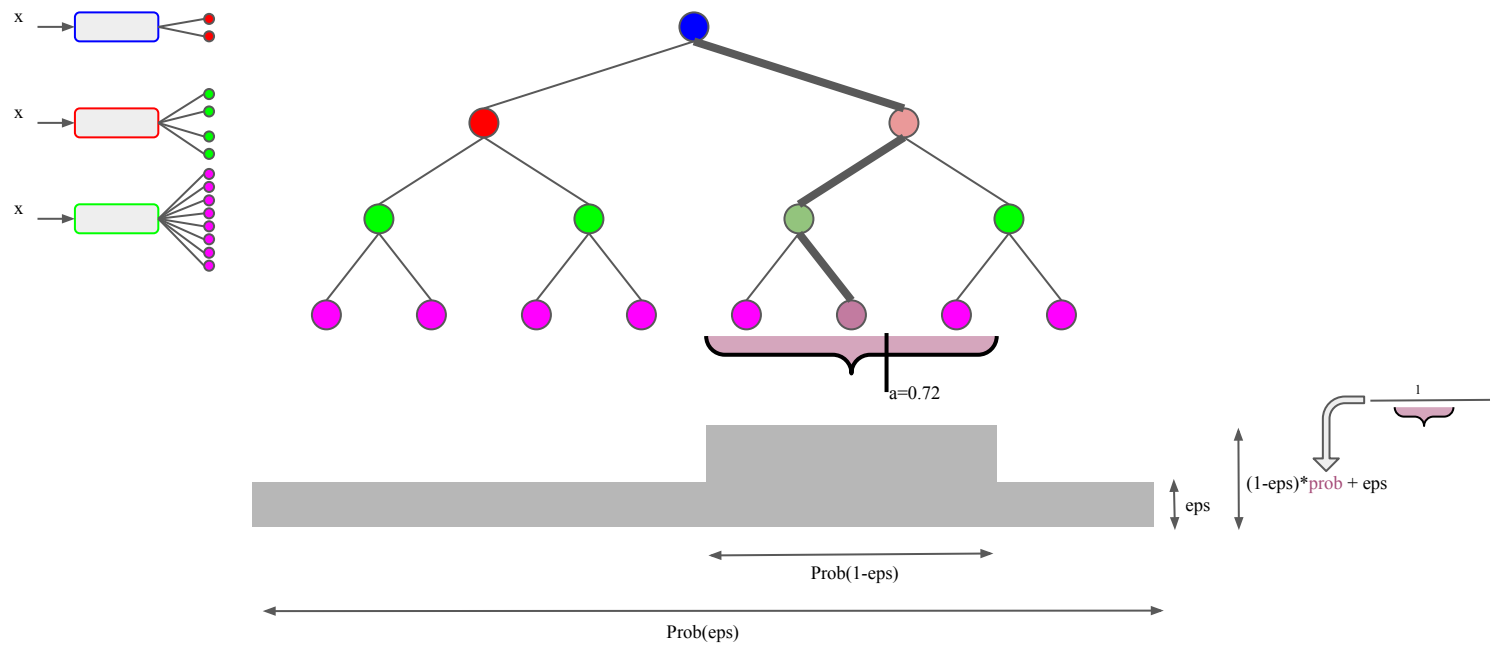
sample an action with eps-greedy



```
def sample(  
    self, obs: Observations, epsilon: float  
) -> Tuple[Actions, Probabilities]:
```

Action query


sample an action with eps-greedy (example: a=0.72)



Action cost

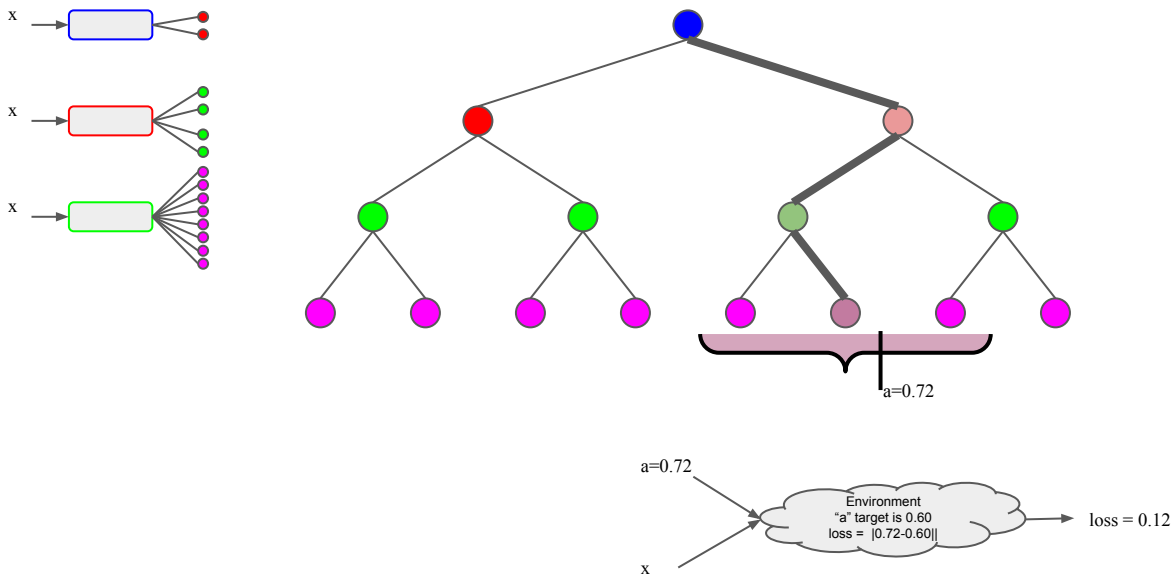
Action cost

apply action in the environment and receive cost feedback



Action cost

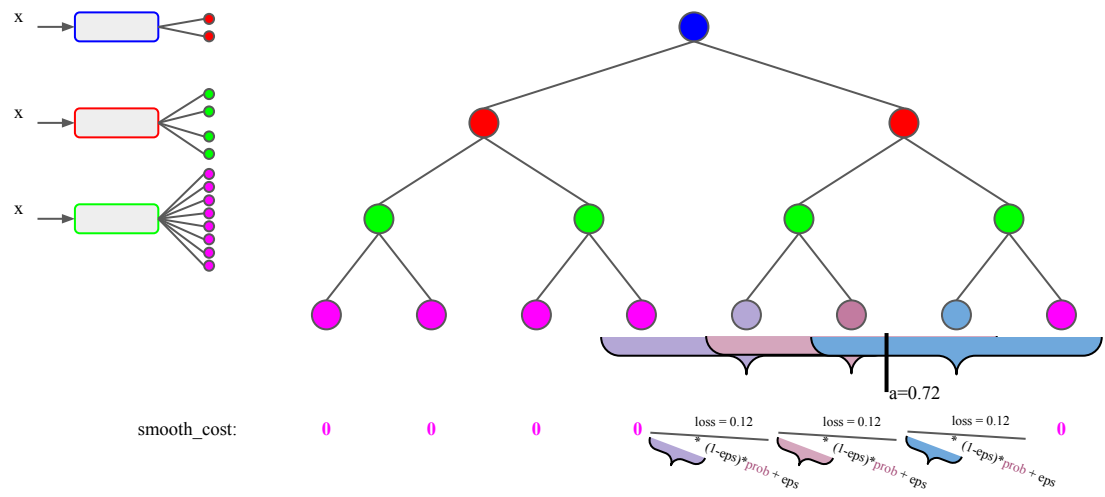
apply action in the environment and receive cost feedback




Action cost

smooth the cost across the discretized actions that could have generated the action

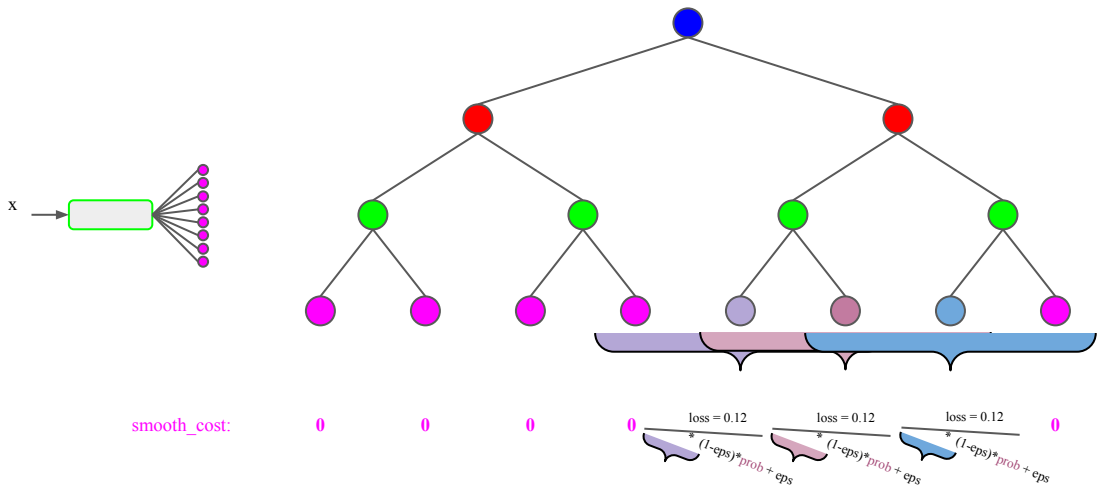
```
@functools.partial(jax.jit, static_argnames=("self",))
def _compute_smooth_costs(
    self, costs: JaxCosts, actions: JaxActions, probabilities: JaxProbabilities
) -> JaxCosts:
```



Update neural network weights

Update: $x \rightarrow$ 


```
def learn(  
    self,  
    obs: Observations,  
    actions: Actions,  
    probabilities: Probabilities,  
    costs: Costs,  
)  
    -> None:
```



$\text{loss ft} = \text{sum}(\text{softmax}(\text{purple circles}) * \text{smooth cost})$

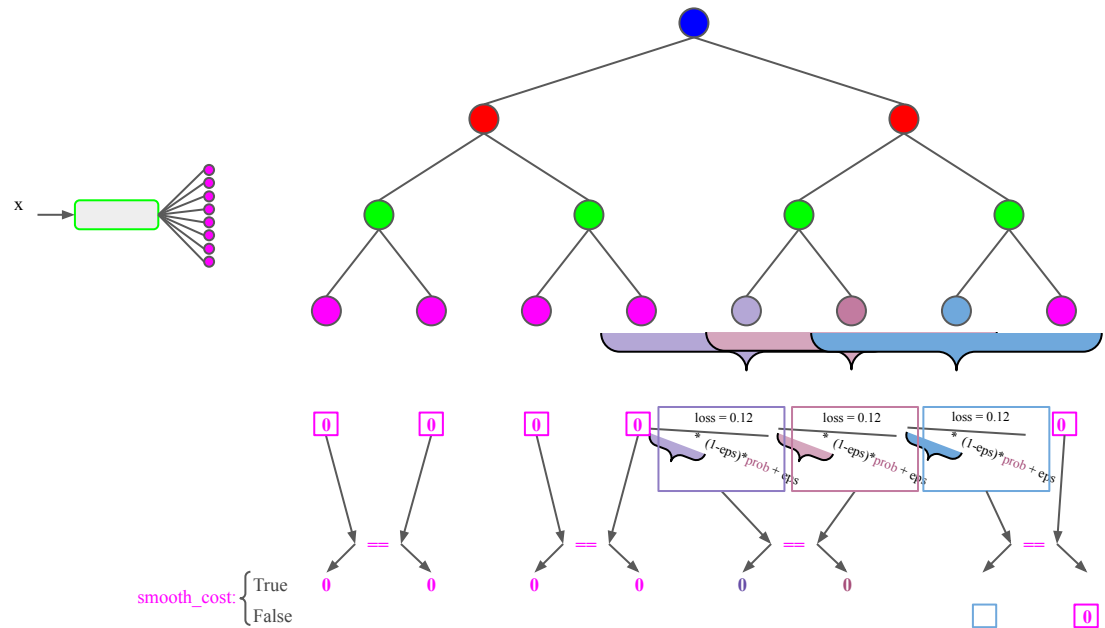
$\text{min}(\text{loss ft})$

```
def learn(
    self,
    obs: Observations,
    actions: Actions,
    probabilities: Probabilities,
    costs: Costs,
) -> None:
```



Update: $x \rightarrow$

Only update nodes whose pair childs have different cost.
In this example:

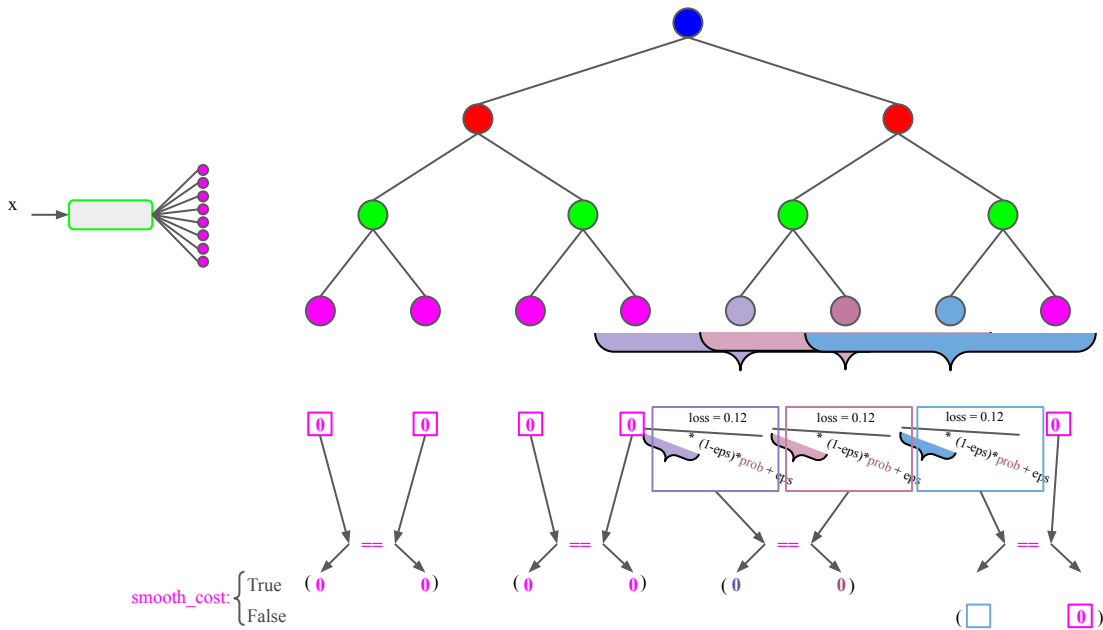


$\text{loss ft} = \text{sum}(\text{softmax}(\begin{smallmatrix} \bullet & \bullet & \bullet & \bullet & \bullet & \bullet & \bullet & \bullet & \bullet & \bullet \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{smallmatrix}) * \text{smooth cost})$
 $\text{min}(\text{loss ft})$

```
def learn(  
    self,  
    obs: Observations,  
    actions: Actions,  
    probabilities: Probabilities,  
    costs: Costs,  
    ) -> None:
```


Update: $x \rightarrow$

Only update nodes whose pair childs have different cost.
In this example:
Note: the softmax is performed pairwise

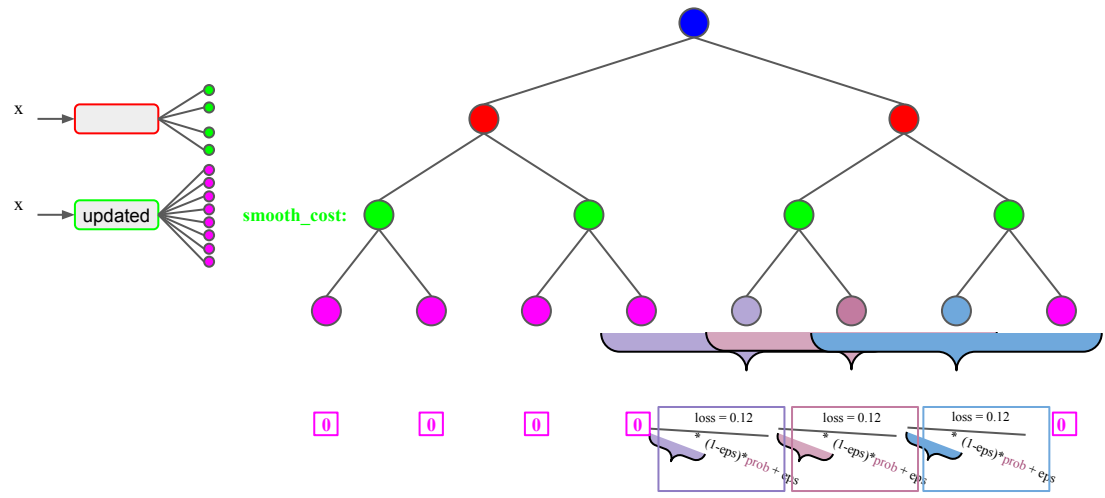


$\text{loss ft} = \text{sum}(\text{softmax}(\text{smooth_cost}) * \text{smooth_cost})$
 $\text{min}(\text{loss ft})$

```
def learn(  
    self,  
    obs: Observations,  
    actions: Actions,  
    probabilities: Probabilities,  
    costs: Costs,  
    ) -> None:
```



Update: $x \rightarrow$ 

```
def learn(  
    self,  
    obs: Observations,  
    actions: Actions,  
    probabilities: Probabilities,  
    costs: Costs,  
)  
    -> None:
```

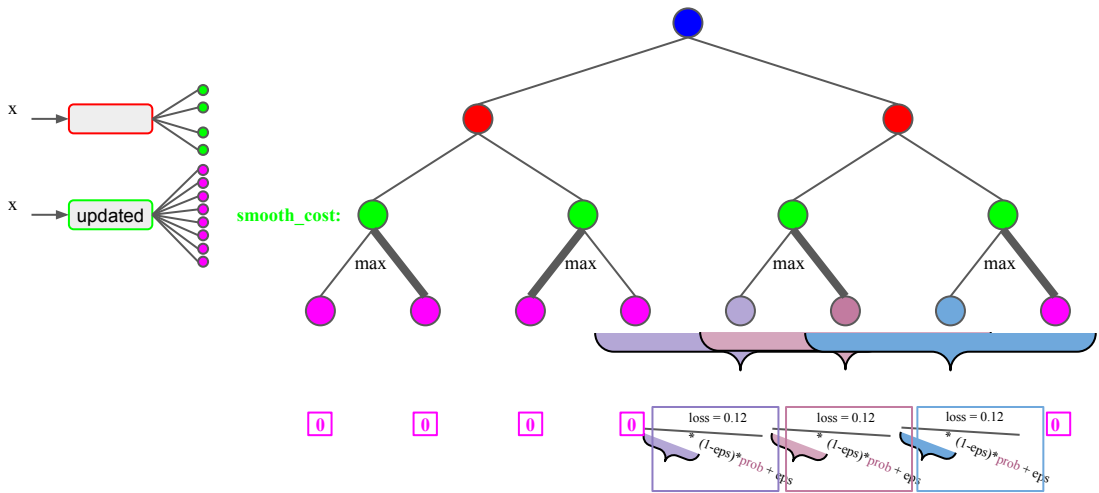




$$\text{loss ft} = \text{sum}(\text{softmax}(\text{green dots}) * \text{smooth cost})$$


$$\text{min}(\text{loss ft})$$

Update: $x \rightarrow$ 

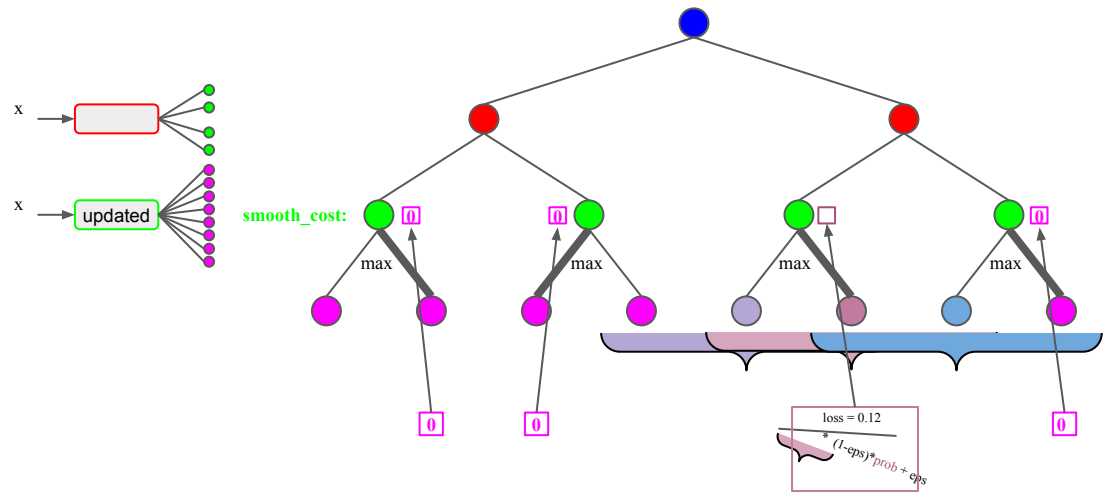
```
def learn(  
    self,  
    obs: Observations,  
    actions: Actions,  
    probabilities: Probabilities,  
    costs: Costs,  
)  
    -> None:
```



loss ft = sum(softmax( ) * **smooth cost**)
min(loss ft)


Update: $x \rightarrow$ 


```
def learn(  
    self,  
    obs: Observations,  
    actions: Actions,  
    probabilities: Probabilities,  
    costs: Costs,  
    ) -> None:
```



$$\text{loss ft} = \text{sum}(\text{softmax}(\text{green_dots}) * \text{smooth cost})$$

$$\text{min}(\text{loss ft})$$


Update: $x \rightarrow$ 




```
def learn(  
    self,  
    obs: Observations,  
    actions: Actions,  
    probabilities: Probabilities,  
    costs: Costs,  
) -> None:
```

$$\text{loss ft} = \text{sum}(\text{softmax}(\text{green circle}, \text{green circle}) * \text{smooth cost})$$


$$\text{min}(\text{loss ft})$$

Update: $x \rightarrow$ 

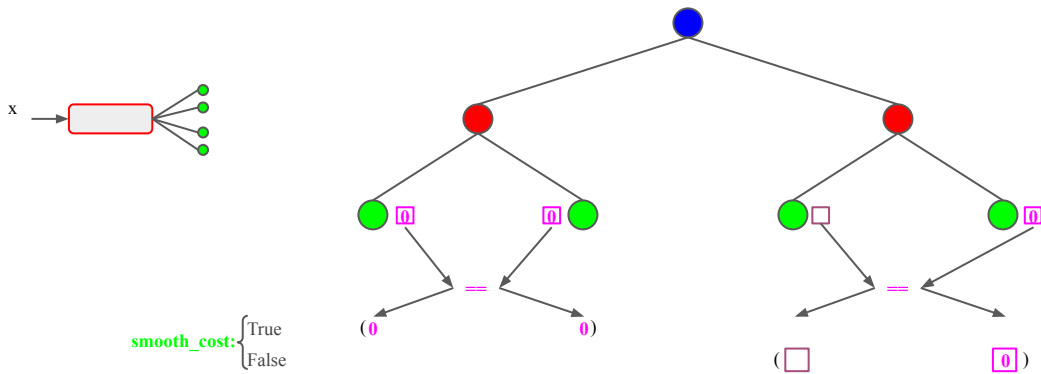
```
def learn(  
    self,  
    obs: Observations,  
    actions: Actions,  
    probabilities: Probabilities,  
    costs: Costs,  
    ) -> None:
```



$loss\ ft = \sum (softmax(\text{green nodes}) * smooth\ cost)$
 $min(loss\ ft)$


Update: $x \rightarrow$ 

```
def learn(  
    self,  
    obs: Observations,  
    actions: Actions,  
    probabilities: Probabilities,  
    costs: Costs,  
)  
    -> None:
```

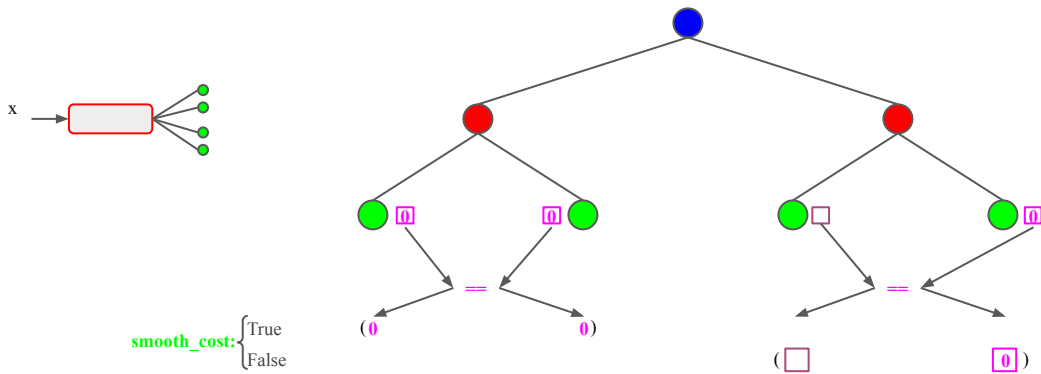


$$\text{loss ft} = \text{sum}(\text{softmax}(\text{green circle}, \text{green circle}) * \text{smooth cost})$$

$$\text{min}(\text{loss ft})$$


Update: $x \rightarrow$ 



```
def learn(  
    self,  
    obs: Observations,  
    actions: Actions,  
    probabilities: Probabilities,  
    costs: Costs,  
    ) -> None:
```



$$\text{loss ft} = \text{sum}(\text{softmax}(\begin{pmatrix} 0 & 0 \\ 0 & 0 \end{pmatrix}) * \text{smooth cost})$$

$$\min(\text{loss ft})$$


Update: $x \rightarrow$ 

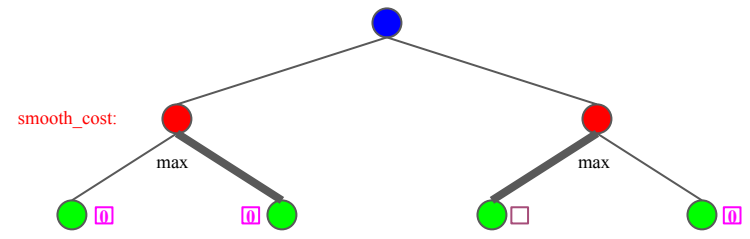
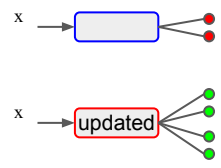


```
def learn(  
    self,  
    obs: Observations,  
    actions: Actions,  
    probabilities: Probabilities,  
    costs: Costs,  
)  
    -> None:
```

$$\text{loss ft} = \text{sum}(\text{softmax}(\bullet \bullet) * \text{smooth cost})$$

$$\text{min}(\text{loss ft})$$


Update: $x \rightarrow$ 


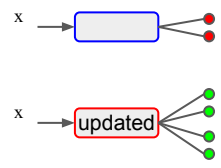


```
def learn(  
    self,  
    obs: Observations,  
    actions: Actions,  
    probabilities: Probabilities,  
    costs: Costs,  
) -> None:
```

$$\text{loss ft} = \text{sum}(\text{softmax}(\bullet \bullet) * \text{smooth cost})$$

$$\text{min}(\text{loss ft})$$


Update: $x \rightarrow$ 




```
def learn(  
    self,  
    obs: Observations,  
    actions: Actions,  
    probabilities: Probabilities,  
    costs: Costs,  
) -> None:
```


$$\text{loss ft} = \text{sum}(\text{softmax}(\bullet \bullet) * \text{smooth cost})$$

$$\text{min}(\text{loss ft})$$

Update: $x \rightarrow$ 




smooth_cost:




```
def learn(  
    self,  
    obs: Observations,  
    actions: Actions,  
    probabilities: Probabilities,  
    costs: Costs,  
) -> None:
```


$$\text{loss ft} = \text{sum}(\text{softmax}(\bullet \bullet) * \text{smooth cost})$$

$$\text{min}(\text{loss ft})$$

Update: $x \rightarrow$ 




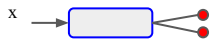
$\text{smooth_cost} : \begin{cases} \text{True} \\ \text{False} \end{cases}$



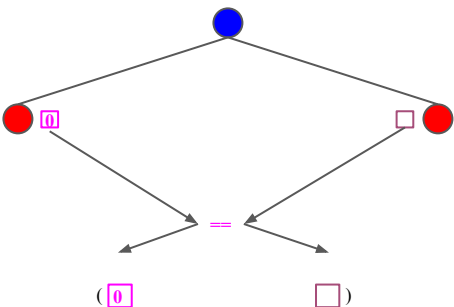
$$\text{loss ft} = \text{sum}(\text{softmax}(\bullet \bullet) * \text{smooth cost})$$
$$\text{min}(\text{loss ft})$$

```
def learn(  
    self,  
    obs: Observations,  
    actions: Actions,  
    probabilities: Probabilities,  
    costs: Costs,  
) -> None:
```

Update: $x \rightarrow$ 



$\text{smooth_cost} : \begin{cases} \text{True} \\ \text{False} \end{cases}$

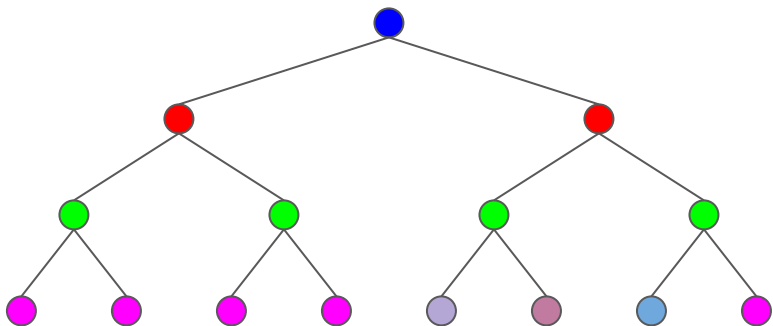
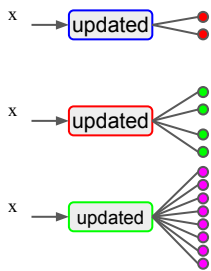


$$\text{loss ft} = \text{sum}(\text{softmax}((\bullet \bullet)) * \text{smooth cost})$$

$$\text{min}(\text{loss ft})$$

```
def learn(  
    self,  
    obs: Observations,  
    actions: Actions,  
    probabilities: Probabilities,  
    costs: Costs,  
) -> None:
```

Update:



```
def learn(  
    self,  
    obs: Observations,  
    actions: Actions,  
    probabilities: Probabilities,  
    costs: Costs,  
) -> None:
```